

Common Steps in GANs (Generative Adversarial Networks):

Training Steps of a GAN (Generative Adversarial Network)

- Step 1: Generate Fake Data
 - The **Generator** starts with random noise and transforms it into fake data that resembles real data.

- Step 2: Discriminator Evaluates Real & Fake Data
 - The **Discriminator** receives:
 - **Real data** (labeled as **1**)
 - **Fake data** (labeled as **0**)
 - It tries to correctly classify whether the input is real or fake.

- Step 3: Compute Loss & Update Weights
 - **Discriminator Loss:** Measures how well it classifies real as **1** and fake as **0**.
 - **Generator Loss:** Measures how well the fake data is classified as **real (1)** by the Discriminator.

- Step 4: Update Weights
 - If the **Discriminator correctly identifies real as 1 and fake as 0**, it **wins**. Generator's weights are updated and discriminator remains unchanged.
 - If the **Generator successfully fools the Discriminator (fake data classified as real)**, the generator **wins**. Discriminator's weights are updated and the generator remains unchanged.
 - **Only one model's weights are updated per training pass.** The losing model is updated. The winning model remains unchanged.

- Step 5: Repeat Until Equilibrium
 - Over multiple iterations, the **Generator improves at creating realistic data**, and the **Discriminator improves at distinguishing real from fake**.
 - At equilibrium:
 - The **Generator's loss is low**, meaning its fake data looks realistic.
 - The **Discriminator's accuracy drops to around 50%**, meaning it can no longer confidently distinguish real from fake.

Key Concepts and Insights on GANs (Generative Adversarial Networks):

Question	Answer
Are GANs supervised or unsupervised?	Unsupervised
Whose weights are updated if Generator wins?	Discriminator
How are weights updated during training?	Alternately
Role of Discriminator during testing?	No role
When does training stop?	At Nash equilibrium
What does low Discriminator loss indicate?	Accurate classification
What does high Discriminator loss indicate?	Realistic data from Generator
Ideal loss value for Generator and Discriminator?	Around 0.5

Callbacks in keras:

Callbacks are functions executed at specific stages of the training process, such as at the start or end of an epoch. They allow monitoring, adjusting learning rates, saving models, or stopping training early.

Purpose:

- Enhance control over the training process..
- Automate tasks like saving the best model or adjusting learning rates.
- Improve model performance and training efficiency.

Example:

- **EarlyStopping:** Stops training when a monitored metric stops improving, preventing overfitting.
- **ModelCheckpoint:** Saves the model at the best epoch based on a monitored metric.
- **ReduceLROnPlateau:** Reduces the learning rate when a metric has stopped improving.

Keras Overview

- **Keras:** High-level API for building neural networks, integrated into TensorFlow.

Common Keras APIs:

API	Description	Example
Sequential API	Linear stack of layers, simple and quick setup.	<code>model = tf.keras.Sequential([...])</code>
Functional API	More flexible, supports complex architectures.	<code>tf.keras.Model(inputs=..., outputs=...)</code>
Model Subclassing	Maximum flexibility using Python inheritance.	<code>class CustomModel(tf.keras.Model): ...</code>

Keras Components:

Component	Description	Examples
Layers	Building blocks of models (Dense, Conv, LSTM).	<code>tf.keras.layers.Dense</code> , <code>tf.keras.layers.Conv2D</code>
Loss Functions	Measures model error (Cross-Entropy, MSE, MAE).	<code>loss='categorical_crossentropy'</code>
Optimizers	Controls learning (SGD, Adam, RMSprop).	<code>optimizer='adam'</code>
Callbacks	Functions triggered during training (e.g., EarlyStopping).	<code>tf.keras.callbacks.ModelCheckpoint</code>

A Simple model Lifecycle:

Question	Answer
Are GANs supervised or unsupervised?	Unsupervised
Whose weights are updated if Generator wins?	Discriminator
How are weights updated during training?	Alternately
Role of Discriminator during testing?	No role
When does training stop?	At Nash equilibrium
What does low Discriminator loss indicate?	Accurate classification
What does high Discriminator loss indicate?	Realistic data from Generator
Ideal loss value for Generator and Discriminator?	Around 0.5